

1 Summary

We seek the solution \mathbf{x} to

$$\mathbf{Ax} = \mathbf{b} \quad (1)$$

where $n \times n$ matrix \mathbf{A} and forcing function \mathbf{b} are known. Our goal is to use row operations to obtain

$$\mathbf{Ux} = \mathbf{d} \quad (2)$$

where \mathbf{U} is an upper-diagonal matrix and \mathbf{d} is related to the forcing function. With an upper-diagonal matrix, the back-substitution algorithm quickly reveals the solution.

2 LU decomposition

2.1 Upper-triangular matrix \mathbf{U}

Beginning with

$$\mathbf{A} \equiv \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} & \dots & a_{1n} \\ a_{21} & a_{22} & a_{23} & a_{24} & a_{25} & \dots & a_{2n} \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} & \dots & a_{3n} \\ a_{41} & a_{42} & a_{43} & a_{44} & a_{45} & \dots & a_{4n} \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} & \dots & a_{5n} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & a_{n3} & a_{n4} & a_{n5} & \dots & a_{nn} \end{bmatrix}$$

we perform a series of row operations to make the first column zero (except for the first entry in the column). This can be achieved via

$$\begin{aligned} (\text{second row}) &\rightarrow (\text{second row}) - \frac{a_{21}}{a_{11}}(\text{first row}) \\ (\text{third row}) &\rightarrow (\text{third row}) - \frac{a_{31}}{a_{11}}(\text{first row}) \\ (\text{fourth row}) &\rightarrow (\text{fourth row}) - \frac{a_{41}}{a_{11}}(\text{first row}) \end{aligned}$$

where the \rightarrow can be interpreted as “becomes”.¹ After proceeding through all n rows, we will have developed an intermediate matrix²

$$\mathbf{A}' \equiv \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} & \dots & a_{1n} \\ 0 & a'_{22} & a'_{23} & a'_{24} & a'_{25} & \dots & a'_{2n} \\ 0 & a'_{32} & a'_{33} & a'_{34} & a'_{35} & \dots & a'_{3n} \\ 0 & a'_{42} & a'_{43} & a'_{44} & a'_{45} & \dots & a'_{4n} \\ 0 & a'_{52} & a'_{53} & a'_{54} & a'_{55} & \dots & a'_{5n} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & a'_{n2} & a'_{n3} & a'_{n4} & a'_{n5} & \dots & a'_{nn} \end{bmatrix}.$$

¹What happens if $a_{11} = 0$, even approximately (e.g., 0.000000017)? We need to rearrange the rows such that this is not the case. Let's not worry about it for now.

²Note the $'$ notation conflicts with MATLAB's implementation of the transpose – this is not a transpose!

If we proceed similarly with

$$(\text{new third row}) \rightarrow (\text{new third row}) - \frac{a'_{32}}{a'_{22}}(\text{new second row})$$

$$(\text{new fourth row}) \rightarrow (\text{new fourth row}) - \frac{a'_{42}}{a'_{22}}(\text{first row})$$

we will arrive at

$$\mathbf{A}'' \equiv \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} & \dots & a_{1n} \\ 0 & a'_{22} & a'_{23} & a'_{24} & a'_{25} & \dots & a'_{2n} \\ 0 & 0 & a''_{33} & a''_{34} & a''_{35} & \dots & a''_{3n} \\ 0 & 0 & a''_{43} & a''_{44} & a''_{45} & \dots & a''_{4n} \\ 0 & 0 & a''_{53} & a''_{54} & a''_{55} & \dots & a''_{5n} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & a''_{n3} & a''_{n4} & a''_{n5} & \dots & a''_{nn} \end{bmatrix}.$$

and can continue with this process until we arrive at

$$\mathbf{U} \equiv \mathbf{A}'''''' \dots \equiv \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} & \dots & a_{1n} \\ 0 & a'_{22} & a'_{23} & a'_{24} & a'_{25} & \dots & a'_{2n} \\ 0 & 0 & a''_{33} & a''_{34} & a''_{35} & \dots & a''_{3n} \\ 0 & 0 & 0 & a'''_{44} & a'''_{45} & \dots & a'''_{4n} \\ 0 & 0 & 0 & 0 & a''''_{55} & \dots & a''''_{5n} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & 0 & 0 & \dots & a''''''_{nn} \end{bmatrix} \equiv \begin{bmatrix} u_{11} & u_{12} & u_{13} & u_{14} & u_{15} & \dots & u_{1n} \\ 0 & u_{22} & u_{23} & u_{24} & u_{25} & \dots & u_{2n} \\ 0 & 0 & u_{33} & u_{34} & u_{35} & \dots & u_{3n} \\ 0 & 0 & 0 & u_{44} & u_{45} & \dots & u_{4n} \\ 0 & 0 & 0 & 0 & u_{55} & \dots & u_{5n} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & 0 & 0 & \dots & u_{nn} \end{bmatrix}.$$

This is the upper-triangular matrix \mathbf{U} that will be useful in the back-substitution algorithm.

2.2 Lower-triangular matrix \mathbf{L}

The row operations described above can be made mathematical by a multiplication by

$$\mathbf{L} \equiv \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & \dots & 0 \\ a_{21}/a_{11} & 1 & 0 & 0 & 0 & \dots & 0 \\ a_{31}/a_{11} & a'_{32}/a'_{22} & 1 & 0 & 0 & \dots & 0 \\ a_{41}/a_{11} & a'_{42}/a'_{22} & a''_{43}/a''_{33} & 1 & 0 & \dots & 0 \\ a_{51}/a_{11} & a'_{52}/a'_{22} & a''_{53}/a''_{33} & a'''_{54}/a'''_{44} & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{n1}/a_{11} & a'_{n2}/a'_{22} & a''_{n3}/a''_{33} & a'''_{n4}/a'''_{44} & a''''_{n5}/a''''_{55} & \dots & 1 \end{bmatrix}$$

as in, $\mathbf{L}\mathbf{U} = \mathbf{A}$. Note that the lower-triangular entries are the same as the multipliers we used in the algorithm to get \mathbf{U} . The creation of \mathbf{L} is practically a by-product of generating \mathbf{U} .

3 The backward-substitution algorithm

Beginning with Eq. (2), we multiply both sides by \mathbf{L} to obtain

$$\mathbf{L}\mathbf{U} = \mathbf{Ld}.$$

Since $\mathbf{LU} = \mathbf{A}$, this reveals that

$$\mathbf{Ld} = \mathbf{b}.$$

Because \mathbf{L} is a lower-triangular matrix, determination of \mathbf{d} is relatively straightforward via the “forward-substitution algorithm”. Framed in terms of a generic lower-triangular matrix with entries ℓ_{ij} , the first three steps are:

$$\begin{aligned}\ell_{11}d_1 &= b_1 \implies d_1 = \frac{1}{\ell_{11}}b_1 \\ \ell_{21}d_1 + \ell_{22}d_2 &= b_2 \implies d_2 = \frac{1}{\ell_{22}}(b_2 - \ell_{21}d_1) \\ \ell_{31}d_1 + \ell_{32}d_2 + \ell_{33}d_3 &= b_3 \implies d_3 = \frac{1}{\ell_{33}}(b_3 - \ell_{31}d_1 - \ell_{32}d_2)\end{aligned}$$

and generally

$$d_i = \frac{1}{\ell_{ii}} \left(b_i - \sum_{j=1}^{i-1} \ell_{ij}d_j \right).$$

We use this to generate \mathbf{d} from the given forcing function \mathbf{b} .

4 The backward-substitution algorithm

Beginning with Eq. (2), and since \mathbf{U} is an upper-triangular matrix, and since we know \mathbf{d} after going through the forward-substitution algorithm above, we can determine the solution \mathbf{x} via the backward-substitution algorithm. Framed in terms of a generic upper-triangular matrix with entries u_{ij} , for a matrix with n rows, the first three steps are:

$$\begin{aligned}u_{nn}x_n &= d_n \implies x_n = d_n/u_{nn} \\ u_{(n-1)(n-1)}x_{n-1} + u_{(n-1)(n)}x_n &= d_{n-1} \implies x_{n-1} = \frac{1}{u_{(n-1)(n-1)}}(d_{n-1} - u_{(n-1)(n)}x_n) \\ u_{(n-2)(n-2)}x_{n-2} + u_{(n-2)(n-1)}x_{n-1} + u_{(n-2)(n)}x_n &= d_{n-2} \implies \\ x_{n-2} &= \frac{1}{u_{(n-2)(n-2)}}(d_{n-2} - u_{(n-2)(n-1)}x_{n-1} - u_{(n-2)(n)}x_n)\end{aligned}$$

and generally

$$x_i = \frac{1}{u_{ii}} \left(d_i - \sum_{j=i+1}^n u_{ij}x_j \right).$$

This expression is mathematically correct (I think), but its implementation is tricky: this algorithm must proceed backwards: $i = n$, then $i = n - 1$, etc.

We use this to get our solution \mathbf{x} .